

All Pairs Shortest Path (APSP) Problem

Given a directed graph $G = (V, E)$, where each edge (v, w) has a nonnegative cost $C[v, w]$, for all pairs of vertices (v, w) find the cost of the lowest cost path from v to w .

- A generalization of the single-source-shortest-path problem.
- Use Dijkstra's algorithm, varying the source node among all the nodes in the graph.

We will consider a slight extension to this problem: find the **lowest cost path** between each pair of vertices.

- We must recover the path itself, and not just the cost of the path.

Floyd's Algorithm

Floyd's algorithm takes as input the cost matrix $C[v, w]$

- $C[v, w] = \infty$ if (v, w) is not in E

It returns as output

- a distance matrix $D[v, w]$ containing the cost of the lowest cost path from v to w
 - initially $D[v, w] = C[v, w]$
- a path matrix P , where $P[v, w]$ holds the intermediate vertex k on the least cost path between v and w that led to the cost stored in $D[v, w]$.

We iterate N times over the matrix D , using k as an index. On the k th iteration, the D matrix contains the solution to the APSP problem, where the paths only use vertices numbered 1 to k .

On the next iteration, we compare the cost of going from i to j using only vertices numbered 1.. k (stored in $D[i, j]$ on the k th iteration) with the cost of using the $k+1$ th vertex as an intermediate step, which is $D[i, k+1]$ (to get from i to $k+1$) plus $D[k+1, j]$ (to get from $k+1$ to j).

If this results in a lower cost path, we remember it.

After N iterations, all possible paths have been examined, so $D[v, w]$ contains the cost of the lowest cost path from v to w using all vertices if necessary.

The Algorithm

```
FloydAPSP (int N, rmatrix &C, rmatrix &D, imatrix &P)
{
    int i, j, k;
    for (i = 0; i < N; i++) {
        for (j = 0; j < N; j++) {
            D[i][j] = C[i][j];
            P[i][j] = -1;
        }
        D[i][i] = 0.0;
    }
    for (k = 0; k < N; k++) {
        for (i = 0; i < N; i++) {
            for (j = 0; j < N; j++) {
                if (D[i][k] + D[k][j] < D[i][j]) {
                    D[i][j] = D[i][k] + D[k][j];
                    P[i][j] = k;
                }
            }
        }
    }
}
```

```

    } } } }
} /* FloydAPSP */

```

Clearly the algorithm is $O(N^3)$.

Finding a Least Cost Path

Floyd's algorithm (modified to find the least cost paths, and not just the cost of the paths) produces a matrix P , which, for each pair of nodes u and v , contains an intermediate node on the least cost path from u to v

So the least cost path from u to v is the least cost path from u to $P[u,v]$, followed by the least cost path from $P[u,v]$ to v .

The following procedure uses the P matrix produced earlier to print the intermediate vertices on the least cost path from node u to node v .

```

Path (int u, int v, imatrix &P)
{
    int k;

    k = P[u][v];
    if (k == -1) return;
    path(u,k);
    cout << k;
    path(k,v);
} /* Path */

```

Note that this procedure could loop forever on an arbitrary matrix, but Floyd's algorithm ensures that we cannot have k on the shortest path from u to v **and** v on the shortest path from u to k .

Proof that Floyd's Algorithm Works

We will prove that after k iterations over the matrix D , $D[i,j]$ is the cost of the cheapest path from i to j that does not include a vertex numbered $> k$.

Proof by induction on k .

Basis: Let $k = 0$ (ie, no iterations yet performed).

Then no intermediate vertices on a path from i to j are allowed, so $D[i,j]$ should be $C[i,j]$ if (i,j) in E , and infinity otherwise. The initialization step does exactly this.

Induction step: Assume that after k iterations, $D[i,j]$ is the cost of the lowest cost path from i to j excluding all vertices from $k+1$ to N .

On the next $(k+1)$ iteration, we are allowed to include vertex $k+1$ in any path.

For all pairs (i,j) , the lowest cost path from i to j excluding vertices $k+2$ thru N goes thru $k+1$ iff there is a low cost path from i to $k+1$ and from $k+1$ to j , excluding vertices $k+2$ thru N .

But the cheapest path from i to $k+1$ without using nodes $k+2$ thru N is simply $D[i,k+1]$ (by the induction hypothesis).

Similarly, the lowest cost path from $k+1$ to j without using nodes $k+2$ thru N is $D[k+1,j]$.

Thus, we should use node $k+1$ to get from i to j iff $D[i,k+1] + D[k+1,j] < D[i,j]$, the cheapest path excluding $k+1$. Since this is exactly what is stored on the $k+1$ th iteration, we have completed the proof.

Comparison with Dijkstra's Algorithm

The all-pairs-shortest-path problem is generalization of the single-source-shortest-path problem, so we can use Floyd's algorithm, or Dijkstra's algorithm (varying the source node over all nodes).

- Floyd's algorithm is $O(N^3)$
- Dijkstra's algorithm with an adjacency matrix is $O(N^2)$, so varying over N source nodes is $O(N^3)$
- Dijkstra's algorithm with adjacency lists is $O(E \log N)$, so varying over N source nodes is $O(N E \log N)$

For large sparse graphs, Dijkstra's algorithm is preferable.